Scala Notes

© Simon Roberts, 2018

Background

Scala is a programming language that can run on a JVM and generate and use class files compatible with Java code, allowing extensive interoperability with Java code and libraries.
Scala supports both object oriented and functional programming paradigms, and has a syntax that borrows elements from the "curly-brace" family of languages. That is, its syntax has similarities with C, C++, Java, JavaScript and others. However, since it adds extensive functional capabilities not provided by those languages, many of the syntax features that support those capabilities are unfamiliar.

Interpreted and compiled modes

- The command `scalac` may be run as a compiler, building binary classfiles, and resulting programs run using the `scala` command

```
$ scalac Hello.scala
$ scala Hello
Hello world!
$
```

- The scala command by itself enters a "REPL" (read, execute, print loop) allowing interactive experimentation

```
$ scala
Welcome to Scala version …
Type in expressions to have them evaluated.
Type :help for more information.

scala>
```

- Scala code can be entered interactively at this point:

```
scala> println("Hello Interactive World!")
Hello Interactive World!

scala>
```

- IntelliJ and other IDEs reflect this interactive mode. IntelliJ provides “worksheets” (Note that IntelliJ uses its own compiler which is not 100% bug free. Worksheets in particular can behave surprisingly and misleadingly in a few situations)
- The interpreted mode can be used to create shell scripts / batch files, though this is inevitably a somewhat operating system-dependent process. In a Unix environment, a script might look like the following. Note that the `#!` characters introduce an interpreter binary path to the shell, and the `!#` marks the end of the preamble for Scala. Everything that follows the `!#` is treated as Scala script:

```
#!/usr/local/scala-2.12.4/bin/scala
!#
val x = 99
println(s"Value is ${x}")
```

Basics

Some immediate comparisons/contrasts with languages like Java, C/C++ are:

- Scala source files which may contain many top-level elements; packages, classes, and objects.
- Source filenames and directories are not required to map to class and package names, though it’s a probably a good

convention to adhere to anyway, since it can assist in finding things during maintenance

- Scala uses packages which can be declared with a statement at the top of the file, or using a block-syntax. For example:

```
package mypackage {
  // package contents
}
```

- An `import` statement can be placed anywhere, and are block scoped, they can also import anything, including object references, fields, and methods. An underscore is used as a wildcard, and braces can group a number of specific items to be imported from a particular source. For example:

```
import math._
import math.{Pi,E}
```

- Many identifiers in Scala appear as though they're part of the core language, but are in fact imports. In Scala an `import` statement can select almost any identifier, even down to fields inside objects. An imported element is available unqualified
- A large number of such identifiers are pre-imported via a standard import of an object called `scala.Predef`. The `Predef` object supplies many features, including `Console.out.println`, which is typically used as simply `println`.
- Note that these seemingly magical imports are common, are frequently provided for parts of the standard API and third party libraries, and can be added to your own libraries.
- Semicolons are not generally required unless multiple statements are on a single line. Because of this, be sure that when wrapping a line, the first and second parts do not form legal syntax individually, or they will be parsed that way. For example:

```
def value = 3
val x = 4
 + value
println("x is " + x)
```

Prints `x is 4`

- Comments are as Java etc.
- The default access in Scala is what public is in other languages. “public” is not a keyword.
- Scala does not have “static” features in a class, nor a static keyword. Instead it uses a “companion object” which is an expressly declared `object` with the same name as a class. The `main` method that serves as the entry point to a Scala program is declared in an object. For example:

```
object MyProgram {
 def main(args: Array[String]): Unit = {
   println("Hello, Scala World!")
 }
}
```

Forms of Identifiers

Scala identifiers take three basic forms

- Letter, letters+digits ($ and _ count as letters,but $ is reserved for compiler-generated ids). Identifiers of this group are broadly familiar to most programming languages including Java. For example:

```
count
_count_of_3
```

  - Identifiers of this form must not be identical to Scala keywords
  - Sequence of “operator characters” (no embedded alphanumerics). Operator characters are printable characters in the 7 bit ASCII set, excluding

alphanumerics, parentheses, brackets, braces, any quotation marks, period, comma, and semicolon. For example:

```
#::
->
^&*%
```

- Almost anything between backquotes, for example:

  `` `😀 smiling happy faces !` ``

  - Note that the backquotes are not part of the identifier
  - This form is primarily intended to allow interoperatbility with Java which exports some identifiers that are Scala keywords, and must be wrapped in this way. For example: `java.lang.Thread.`yield``

Types, Literals

- Scala types all mimic pure objects (though the underlying implementation might use primitives for efficiency)
- For example, basic types include: `Boolean`, `Int`, `Float`, `String`
  - Note that the types that map to primitives have capital first letters
- The root type in Scala's hierarchy is `Any`, this has two immediate subtypes `AnyRef` and `AnyVal`,
  - `Any` serves as the most general type (much as `Object` does in Java, except it can refer to the primitive-like types too)
  - `AnyRef` maps directly to `java.lang.Object`
  - `AnyVal` is the parent class of the primitive-like types. In addition Scala allows user defined value types to be defined subclassing this. In such a situation, the class

must define exactly one field of primitive type. The type is named, is not the same type as the primitive it is built on (which can be useful for type safety), and can contain methods as normal.

- `Nothing` is a special type that is assignment compatible to all types (it's the "bottom" of the type hierarchy, a subtype of every other type). It can be used to represent the absence of a value where one might be expected (for example a method that throws an exception instead of returning a value.)
- The empty pointer value is `null` , the type of this is `Null`
- Literal forms for numbers are largely as expected. As with Java, default types are `Int` for integral types, and `Double` for floating point types. For example:

```
123
1.23
1.23E+4
1234567890123L
1.23F
```

- Notes:
  - Literals of integral types must not begin with redundant zeroes
  - Binary and octal forms are not supported
  - Underscore grouping separators are not supported
- Simple character and String literal are also familiar, a character is surrounded with single-quotes, Strings with double-quotes. These literals use backslashes for escapes and Unicode. For example:

```
'X'
'\u96E5'
"Hello\n"
"Hello\\not a newline"
```

- A String literal can be created such that special character handling is disabled using the prefix raw immediately prior to the opening double quote mark
- `raw”Hello\not a newline”`
- Or using three double quotation marks. In this case, a string literal can contain newline characters directly:

```
"""
Hello there
Line 2!"""
```

- Note that if the string is indented for formatting in the source, those indent characters will also form part of the string. This can be avoided using the vertical bar character and the “stripMargin” operation:

```
"""Hello
  |Line 2
""".stripMargin
```

- Scala provides range types, with literals using the words “to” (denoting an inclusive range) and “until” (denoting a range that excludes the upper bound). A step size may be specified using the word “by”. For example:

```
1 to 10
1 until 10
1 to 96 by 5
1 to -10 by -2
1 until 96 by 5
```

Variable Declarations, Type Specification, and Inferencing

Three main forms of “variable” declaration are provided, `val`, `def`, and `var`

- `val` associates an identifier with a value, which must be assigned exactly once
    - This is not a “constant” if the object referred to is mutable
- `def` declares an association between an identifier and some means of obtaining a value. This might be a function/method, or a value. This can also be an abstract association that is satisfied in subclassing (by either a `val` or a `var`)
- `var` associates an identifier with a mutable value. (Note that “good Scala style” encourages functional approaches and immutable values.)
- In interpreted mode a `val` might seem to be mutable, as the interpreter permits re-declaring the identifier.
- Declarations use `val`, `def`, or `var` to introduce the declaration, followed by the identifier.
- The type may be specified with a colon followed by the typename.
- An assignment that initializes the value (or provides the body of a function) may follow.
- Where a variable is initialized, the type may often be inferred and omitted from the source.
- Example:

```
val count = 3 // inferred to be Int type
val count: Int = 3 // explicitly Int type
```

- A val declaration may not be changed after assignment
- Variable declarations are one per statement, the comma separated list of declarations that’s common in other languages is not supported:

~~`var x = 2, y = "Hello"`~~ `//` ***NOT LEGAL IN SCALA***

Expressions

- Typical expression forms familiar from C/C++/Java using infix operators work as expected with type inferencing/promotions that are unsurprising. For example:

```
val x = 2 + 9.2 // x is a Double, value 11.2
```

- The plus operator '+' performs String concatenation and conversions of non String operand types in mixed-type operations
- Less familiar is that *blocks* also generally have an expression value. They adopt the value of the last expression in the block. For example:

```
val b = { println("Hello"); 7 } - 3 // b is 4
```

- The `if` construction is a conditional expressions, similar to the ternary operator in C/C++/Java etc. As before, blocks have value. For example:

```
val message = "There are " +
  (if (count > 3) "many" else "a few")
  + " items"
```

  - As with other languages, it's often advisable to use braces for the `if` and `else` subordinate expressions if they are placed on separate lines
  - As with Java, but different from C/C++/JavaScript, the expression to be tested must be a `Boolean` expression.
- Quoted strings can be variable expressions, evaluated at runtime, that describe substitutions (known as interpolation) based on expressions embedded in the string. Two forms exist, defined by prefix letters placed immediately prior to the opening double-quote mark. The letter may be `s`, or `f` (both

are lowercase). The letter `s` performs a simple expression substitution, whereas using `f` allows a formatted substitution comparable to a `printf` type effect. The expression is prefixed with a dollar symbol and enclosed in braces (like a Unix shell variable) if it is other than a simple identifier. For example:

```
val x = 9; val y = 3
println(s"$x is divisible by $y " +
        s" is ${x % y == 0}")
println(f"3--more precisely ${math.Pi}%7.2f")
```

- Operators in Scala are actually methods on objects
- All methods can be invoked in an infix syntax. For a single argument method this translates the first form directly to the second form shown here:

```
val s = "Hello, World!"
val pos = s.indexOf('W') // pos is 7
val pos2 = s indexOf 'W' // equivalent
```

- Scala allows programmer defined operators and operator overloading. Because arbitrary new operators may be invented, the precedence rules are a little more complex than in other languages. Generally, operator precedence is based on the first character of the operator name.

Highest precedence

Special characters not otherwise mentioned

```
* / %
+ -
:
< >
= !
&
```

^

|

All letters

Lowest precedence

- An exception to this is that assignment operators (`+=` etc.) take very low priority

Arrays, Lists and Tuples

Arrays:

- Carry their types as a generic type argument
- May be created with a capacity using `new Array[Int](count)`
- May be created with contents using `Array[String](item1, item2)`
- Elements are accessed at subscripts using parentheses, not square brackets
- Array size may be determined using `length` or `size` element

```
val empty: Array[Int] = new Array[Int](5)
println("size is " + empty.size) // 5
println("first element is 0 is " + empty(0))
empty(0) = 99
println("first element is 0 is " + empty(0))

val names: Array[String] = Array("Fred", "Jim",
"Sheila")
```

- Reading an array uses the `apply(index : Int)` method behind the scenes, this is why Scala uses parentheses, not square brackets for subscript access
- Assignment to an array element uses the method

```
update(index: Int, x: Element)
```

- Lists
  - Generally immutable, although mutable variants exist
  - May be constructed using

```
List("Fred", "Sheila")
```

  - Are usually accessed via their `head` element and “the rest” via `tail`
  - Item-at-index is extracted as with arrays
  - Changes result in a new List

```
val ln = List("Harry", "Sally", "Pally")
println("List is empty? " + ln.isEmpty) // false
println("List contains " + ln.length) // 2
println("Head of list is " + ln.head) // "Harry"
println("Tail of list is " + ln.tail)
// List("Sally", "Pally")
```

- Tuples allow for grouping multiple items in one bucket without declaring and instantiating a special class and object.
  - Tuples use a simple parentheses type syntax, which can be on the left side of an assignment.
  - Tuple structures may appear on the left side of an assignment which results in assignment of all the named fields at once. (Although perhaps confusing, this can work with assignment of actual parameters to a single formal parameter of tuple type in a function call).
  - Tuple elements are accessed with a dot syntax, using names `_1`, `_2` etc. The numbers indicate the position of the item being addressed.

```
val t1 = ("Fred", 12, true)
println(t1._1) // "Fred"
val (a, b, c) = (9, math.Pi, "Ostrich")
println(s"values are a = $a, b = $b, c = $c")
```

- Tuples are instances of `Product<n>` traits, these allow:
- Determining the number of elements
- Accessing an element by its computed position
- Obtaining an `Iterator` over the elements

Defining classes

Scala fully supports the notion of classes, which may be defined in hierarchies, as abstract or concrete, and in a special form called a “case class”.

- A simple class with no fields or behaviors may be declared:

```
class VerySimple
```

- Note that classes, fields, and methods are “public” by default and Scala does not have a `public` keyword

Constructors

- The body of the class amounts to the primary constructor, code within it is executed from top to bottom (defining methods and fields as it goes).
    - Arguments to the “primary” constructor are declared after the classname
    - Constructor accessibility (which can be reduced) prefixes the argument list

```
class VerySimple private(x: Int) {
  println(s"Running constructor with arg $x")
  val myX = x
}
```

- If a constructor argument is to be used directly as a field, it can be declared as such in the constructor argument list by adding either a `val` or `var` prefix:

```
class VerySimple(val x: Int) { …
```

Auxiliary constructors

- Auxiliary constructors are overloaded constructors taking differing argument lists. They must invoke one other constructor first using a `this(arguments)` syntax.
- Accessibility of main and auxiliary constructors may be specified independently

```
class VerySimple private(val x : Int) {
  def this(x:Int, s:String) = { // not private
    this(x) // delegate to primary constructor
    println(s"Auxiliary constructor s = $s")
  }
  println(s"Running constructor with arg $x")
…
```

Methods / Functions

- The basic form of a function or method declaration is:

```
def name(arg: ArgType, arg2: ArgType) :
ReturnType = body-expression
```

- The parentheses and argument list may be omitted if no arguments are expected.
  - In this case, the invocation *must not* use parentheses either
  - If empty parentheses are used on the declaration, parentheses are optional on the invocation
  - By convention, do use parentheses if your method mutates internal state or has side-effects. Omit them if the method is a pure function.
- The return type may be omitted if:

    - The return type can be inferred from the `body-expression`, and
    - The function is *not* recursive
    - It's considered poor style to omit the return type from public methods
- The `body-expression` is an expression that computes the value to be returned by the function call
    - Commonly, this is presented as a block (enclosed in curly braces), but any single expression is legal
    - The expression value of a block is the last expression in the block
- In Scala the `return` keyword is generally *not* used (but it exists and can be used)
- Functions may be declared in objects, classes, and functions. Those declared in objects and classes are generally referred to as methods
- Scala does not have "static methods" instead such behaviors may be declared in a companion object
- Methods can use operator-like identifiers for their names
- Method overloading is supported, but the keyword `overload` is required
- Scala uses the `this` keyword similarly to other languages as a reference to the object on which a method was invoked

Example methods:

```
// explicit return type
def sayHello(name : String) : Unit =
  println(s"Hello $name")

// return type inferred
def makeGreet(name: String, handle: String) =
  handle + " " + name
```

```
// block expression
def makeGreet(name: String, isFem: Boolean) = {
  val handle = if (isFem) "Ms." else "Mr."
  handle + name // "return value"
}

// override
override def toString: String =
  s"VerySimple, x is $x"
```

- A method named `apply` has special use. If defined, it is invoked by placing parentheses directly after the reference to the object to which the method belongs.
- Similarly, a method named `update` is used in assignment situations:

```
class ArrayLike {
  def apply(x: Int): Unit =
    println(s"called apply with argument $x")
  def update(idx: Int, x: Int): Unit =
    println(s"updating index $idx to value $x")
}

val e = new ArrayLike
e(33) // "called apply with argument 33"
e(33) = 101 // "updating index 33 to value 101"
```

Variable Length Argument Lists

Variable length argument lists can be defined

- The variable part must be the last item in the list
- Indicate that it's variable by following the type with an asterisk
- The received argument will be some kind of sequence, the precise type depends on the invocation

- Simple invocation requires only comma separated arguments
- Note that the single “fixed” argument in the example is illustrative, not required

```
def addUp(y: Int, x:Int *): Int = x.sum + y
println(addUp(1, 2, 3)) // prints 6
```

- Variable length argument lists may be supplied in calls by sequence like structures
  - This is done using the _* syntax
  - This is likely to result in a different type for the receiver, but it will still be sequence-like

```
println(addUp(1, List(2, 3): _*))
```

- Note that variable length argument lists cannot take default values

Default Arguments

- Default argument values may be specified using an assignment, the equals sign comes after the argument’s type specification. The type specification is not optional even if type could be inferred from the default value:

```
def makeGreet(name:String, isFem:Boolean = true)
```

Recursive functions and tail-recursion

- Recursion is the basic mechanism for iteration in a pure functional language, as such it’s used quite frequently in Scala code
- If a function is recursive, it must explicitly declare its return type
- A recursive function might use a lot of stack space, possibly overflowing and crashing the thread. This might be avoided if the function can be made “tail recursive”

- In a tail recursive function the recursive invocation is the very last thing performed by the function.
    - Scala compiler will “unwind” a tail-recursive function call and convert it into a loop
    - Tail recursion can be more efficient in execution time
    - Tail recursion will not overflow the stack no matter how many iterations are needed
    - If a function is intended to be tail-recursive, adding the annotation `@tailrec` will give an error if the compiler determines that it is actually not
    - Converting a function to a tail-recursive implementation often involves adding an accumulator of some kind to pass intermediate results into the recursive call, rather than accumulating the result *after* the recursive call returns

```
// NOT tail recursive, annotation would fail
def len(l:List[Int]): Int = l match {
  case List() => 0
  case h :: t => 1 + len(t)
}

// tail recursive nested function
def trLen(l:List[Int]): Int = {
  @tailrec def accumulateLen(acc: Int, l:
List[Int]): Int = l match {
    case List() => acc
    case h :: t => accumulateLen(acc + 1, t)
  }
  accumulateLen(0, l)
}
```

Pass By Name Arguments

- Arguments passed to a function invocation are generally evaluated before the method invocation. However, using a “by name” parameter, evaluation can be delayed. This can be valuable if for example the argument might never be used, or to allow exception handling to be performed inside the method rather than at the call site:

```
// msg argument passed “by name”
// not evaluated at invocation time
def log(level: Int, msg: => String) : Unit = …
```

Currying

Where a function takes multiple arguments, these are conventionally defined in a single argument list, but may alternatively be presented as multiple argument lists. This is called “currying”. In this case, it’s usual to put a single argument in each list.

- These two functions are broadly equivalent:

```
def add(a:Int, b:Int): Int = a + b
def addCurried(a:Int)(b:Int): Int = a + b
```

- Simple invocation follows the parentheses of the definition:

```
add(2, 2)
addCurried(2)(2)
```

- Note that these particular two functions cannot have the same name, because the type and order of arguments is the same, they are not distinct overloads
- A key use of currying is to support “partial invocation”
  - Partial invocation creates a new function that takes fewer arguments than the original
  - The “missing” arguments have been pre-defined

- In this example, the type of addTwo is expressly declared as "function taking single `Int`, returning `Int`"
- In this example, the expressly defined type is required, or the compiler will object to "`addCurried(2)`" stating that more arguments are needed

```
def addTwo: Int => Int = addCurried(2)
println(addTwo(4)) // prints 6
```

- In the absence of expressly defined type in the context, appending an underscore to the partially applied form provides an explicit way of requesting the curried function.

```
def addTwo = addCurried(2)_
```

- From a formal perspective, a curried function of the form

```
def addCurried(a:Int)(b:Int): Int = a + b
```

  Is a function that takes an `Int`, and produces a function that takes an `Int` and produces an `Int`, like the following example to which it is functionally identical:

```
def addCurried: Int => Int => Int =
    (a:Int) => (b:Int) => a + b
```

- Note that if you "manually curry" your functions, you will not need to use the underscore syntax to extract the partially applied function:

```
def addTwo = addCurr(2)_ // for Scala currying
def addTwo = addCurr(2)  // for manual currying
// Note the underscore ^ or absence of it
```

- Curried functions are often declared to assist with the type inferencing system.
    - Inferencing works on each argument list as a whole
    - Type information is passed from left to right, so it's often possible that a curried function works better with behavioral arguments (lambda expressions) than a regular one.
    - This is particularly effective if the lambda expressions are used in argument lists at the right of the group

Braces invocation

- Where a function takes a single argument in a set of parentheses, curly braces may be substituted for the parentheses
- This can create a kind of "psuedo-syntax"--that is a function invocation that looks more like a new language feature
    - This style is particularly common when using curried functions

```
def repeat(count: Int)(op: => Unit): Unit = {
  for (i <- 0 until count) op
}

repeat(3) {
  println("hello")
}
```

Infix invocation

Methods, regardless of the identifier used, can be invoked using a traditional dotted syntax or an infix form

- Infix notation mandates that the `this` object goes before the method name
- Infix notation is not practical where `this` is not available, such as for functions nested inside other functions
- Infix notation can be used when more than one argument is required; place the arguments in parentheses like a normal argument list

```
// given:
class Value(val x : Int) {
  def times(y:Int) = x * y
  def twoArgs(y:Int, z:Int): Int = x + y + z
}
val v = new Value(2)
// infix invocations
println("v times 3 is " + (v times 3))
println("v twoArgs (3, 4) is " +
  (v twoArgs (3,4)))
```

Defining Subclasses

- A class may be defined as a subclass of one other class using the `extends` keyword
- Scala permits only a single class parent, but traits support multiple inheritance of both state and behavior
- Subclassing may be prevented by using the modifier `final` on the class
- A class definition may carry the keyword `sealed`, in which case all subclasses must be defined in the same source file. This permits some useful optimizations.
- The parameter list for the parent constructor is appended to the parent classname in the extends clause
- The parent may be initialized using any of its constructors, either primary or auxiliary

- Access to parent features is supported using the super keyword
- Parent features may be protected or private, but only protected features may be used

```
sealed class Base private (x: Int) {
  def this(x: Int, b: Boolean) = {
    this(x)
    println("In auxiliary base constructor")
  }
  override def toString: String =
    s"Base, x = $x"
}

final class Child(x: Int, y: Int) extends
 Base(x + 7, true) {
  override def toString: String =
    s"Child, y = $y, from ${super.toString}"
}
```

Abstract classes

- A class may be declared `abstract`, in this case `val`, `var`, and `def` elements may be left uninitialized (but must be explicitly typed). Any such must be satisfied with a real implementation in any non-abstract derived class.
    - Using `def` is generally preferred
- An abstract `val` element in a base class requires a concrete `val` implementation
- An abstract var element in a base class
- An abstract `def` element may be implemented using `def`, `var`, or `val`
- An abstract var element may be implemented by a var or a pair of def methods, one a setter, and the other a getter.

    - Getter is named the same as the var
    - Setter is varname with _= appended
    - The getter/setter approach allows, for example, an easy transition from a simple variable to a range-checked assignment while retaining simple assignment source semantics.

```
abstract class Base {
 def a : Int
 var p : Int
}
class Child(x: Int) extends Base {
 val a = 9 // def implemented by val
 def p = 3 // val implemented by get/set pair
 def p_= (x: Int) = println(s"Setting p to $x")
 override def toString: String =
   s"Child, x = $x, a = $a, p is $p"
}

val c = new Child(3)
println(c.p) // prints 3
c.p = 99     // assignment prints 99
```

Traits

Scala provides a mechanism for multiple implementation inheritance, which is the “trait”.

- Traits are similar to abstract classes
- They cannot be instantiated directly, but must be used in the manner of parent classes for creating other objects
    - Those objects can be anonymous subclasses of the trait
- They can have both concrete and abstract fields and behaviors

    - Identically named accessible features from multiple parents will cause compilation errors
- Traits cannot have constructor arguments
- A single class can can be assembled with multiple traits
    - If the class also has a parent class, use this form for the declaration:

```
class MyClass extends MyParentClass with
  TraitOne with TraitTwo
```

    - If the class only inherits from traits use this form

```
class MyClass extends TraitOne with TraitTwo
```

- A trait may demand a particular parent (possibly anonymous) class if needed to satisfy its general behavior

```
trait T extends A …
```

    - Features of the required parent may be overridden and may be accessed using `super`
    - Where multiple traits are listed, this can provide a decorator-like chainable behavior--the behaviors are invoked from right to left (based on the `with X with Y` declaration)
    - The required parent may be abstract in which case the concrete parent implementation (and with it the meaning of `super`) is supplied when the trait is mixed into to some other class
    - If the parent type is abstract, overriding methods must be marked `abstract override`

```
abstract class A { def show: Unit }
trait T extends A {
  abstract override def show: Unit = {
    print(" T"); super.show; print(" xT")
  }
}
```

```
trait U extends A {
  abstract override def show: Unit = {
    print(" U"); super.show; print(" xU")
  }
}
class C extends A {def show:Unit = print(" C") }

new C with T show; // T C xT
new C with T with U show; // U T C xT xU
new C with U with T show; // T U C xU xT
```

- This approach is very flexible, allowing creation of anonymous classes that are made from any combination of base class plus mixed-in traits, however, the traditional OO “decorator” pattern provides additional flexibility in that the order and number of decorating elements may be changed at runtime.
    - If dynamic changes are definitely not required, then the trait mix in syntax is significantly more compact

Anonymous classes

- Anonymous classes may be derived and instantiated
    - This may be done from concrete or abstract base classes, from traits, or without any explicit parent

```
abstract class GetDoubled {
  def v: Int
  def get: Int = v * 2
}

val v = new GetDoubled{val v = 2};
println(v.get) // prints 4
```

- Without any parent::

```
new {def show: Unit = println("Hello!")}.show
```

Objects (Singletons) and Companion Objects

Scala does not provide a keyword “static”, instead the object keyword declares and instantiates a single instance.

- The result is an instance of an object (the defining class name is related to the object name, but is not the same)
- The object behaves as an object and elements may be addressed using **`this`** if desired

```
object OnlyOne {
  val count = 100
  val message = "Hello"
  def main(args: Array[String]): Unit = {
    println(s"count is $count, " +
      s"message is ${this.message}")
  }
}
```

- Features that relate to a concept as a whole, rather than an individual instance of a class can be implemented in a “companion object”, which is an object with the same declared name as the class.
    - Even though the object’s class is actually different, it shares access to private elements
    - Defining an apply method on the companion object is a common pattern for instantiating objects without the keyword new in client code:

```
class MyClass private (val x : Int) {
  override def toString: String =
    s"instance of MyClass, x is $x"
}
```

```
object MyClass {
  def apply(x: Int): MyClass = new MyClass(x)
}

val mc = MyClass(3)
println(mc)
```

Enumerations

Scala does not have a language element to create values selected from a pre-determined set of legitimate values, but instead uses a library class as the basis of enumerations. This is the `Enumeration` class.

To define a simple enumeration, create an *`object`* that extends the class `Enumeration`, and create the required values by declaring the fields, and initializing them using the `Value` method. To provide an identifiable type for the enumeration, capture the type of the `Value` method in a type definition that's enclosed in the enumeration object:

```
object DayOfWeek extends Enumeration {
  type DayOfWeek = Value

  val SATURDAY = Value("Saturday")
  val SUNDAY = Value("Sunday")
  //... more repetitions omitted
}
```

GIven this declaration and appropriate imports, other code can refer to individual values such as: `DayOfWeek.SUNDAY`, and can

declare variables such as:

```
val today: DayOfWeek.DayOfWeek =
  DayOfWeek.MONDAY
```

The `match / case` Construct

Scala provides a very powerful and flexible construction that is somewhat similar to a case statement in other C-like languages.

- In its simplest form it can match literals

```
val s = "bonjour"
s match {
  case "bonjour" => println("Hello!")
  case "au revoir" => println("Goodbye!")
}
```

- It can also match a default-like situation, capturing the value

```
case x => println(s"What does $x mean?")
```

- If the default value is not needed in the right hand side, an underscore (Scala's general wildcard) may be used instead

```
case _ => println("No idea what that meant!")
```

- Matches might overlap, and are tested from top to bottom. The first match wins, and no others are tried.

- If a match clause exhausts all cases without matching, an exception is thrown
  - `case x =>` … may be added at the end to pick up unmatched situations

- Matches can match by type, which is generally much preferred to using `target.isInstanceOf[Type]` tests

```
val x:Any = s
x match {
  case i:Int => println(s"int value $i")
  case s:String => println(s"string $s")
  case _ => println("Some other type")
}
```

- Matches can be “guarded” with if clauses, which can benefit from the matched type (note that parentheses are not actually required around the boolean parameter to the `if` clause, even though normal `if` statements do require them).

```
case s:String if (s.length > 3) =>
    println("A string longer than 3")
```

- For many types (e.g. `List`) a match can extract how the object was created, and/or decompose it. Note that the `::` (“cons”) operator describes adding a new head (the left operand) to an existing list (the right operand). In a match statement, the joined parts are extracted

```
val l = List(1, 2, 3)
l match {
 case head :: tail => println(s"head is $head")
}
// These also match:
case List(x, _, _) =>
  println(s"List with 3 elements, first is $x")
case List(x, _, _*) =>
  println(s"List with 2+ elements, first is $x")
```

- When matches are made with patterns that extract elements form the item being matched, it’s also possible to capture the entire matched item using a “pattern binder”

```
case l @ List(x, _, _) =>
```

```
  println(s"List with 3 elements: ${l}")
```

this form would have the entire list printed from the variable `l`.

- A match can be made against a variable. This requires a special syntax. Recall that this form:

```
val x = … // numeric value
x match {
  case i => println(s"int value $i")
```

"captures" the value that is being matched in the variable `i`. If a variable `i` already exists, this will create a more local variable that *shadows* the original. Instead, this syntax:

```
val x = … // numeric value
val i = … // numeric value
x match {
  case `i` => println(s"x matched $i")
```

does not create a new variable i, and matches when the values of x and i are equal.

Implementing `FunctionN` and `PartialFunction`

- If a context (e.g. function actual parameter) requires a `Function` or `PartialFunction`, a match block can be provided directly, without needing to specify the "`x match`" prefix. For example, these two are equivalent:

```
val f2: Function2[Int, Int, String] = {
  case (x, y) if x + y == 5 => "Five"
}

val f3: Function2[Int, Int, String] =
```

```
  (x,y) => (x,y) match {
    case (a, b) if a + b == 5 => "Five"
}
```

Programmer-defined match targets

Matching against class type can be used to define and group variant behavior, achieving a similar goal to method overriding. In this case however the variations are placed in a “manager” code block, unrelated to the particular classes. This *might* be a cleaner, lower maintenance, design in some circumstances, such as when combinations of related types interact in ways that are properly managed from *outside* the instances. In such a situation, one often finds a conventional OO approach leads to a combinatorial explosion of “classification” and “action” behaviors for every class. This is often better handled on the outside and a `match` statement can be very elegant for this.

- Case classes allow easy definition of classes that are optimized for use with matching

```
sealed class Transporter()
sealed case class Car(seats:Int) extends
  Transporter
sealed case class BoxVan(length:Int, width:Int,
height:Int) extends Transporter

val t: Transporter = new BoxVan(3,4,5)
t match {
  case Car(s) =>
    println(s"Item on seat 2 of $s")
  case BoxVan(l, w, h) =>
    println(s"van capacity ${l*w*h}")
}
```

- Scala also allows any class to be built in a way that integrates with matching by providing an “extractor” behavior.

Handling Exceptions

- Scala uses an exception mechanism broadly similar to other languages
- Exceptions are objects, and they have type which can be used to control how they are handled
- Exceptions may be thrown by any code, and that code does not have to declare that it does so. That is, Scala exceptions differ from Java in that there is no “declare or handle” rule and no concept of checked exceptions.
- Handling exceptions is performed with a `try catch finally` structure.
- The single `catch` block is not type-specific, does not define a formal parameter, but uses a match construction to partition the various catch behaviors

```
 def mightWork:Unit = {
    val result = math.random
    if (result < 0.34)
      throw new SQLException("DB broke!")
    else if (math.random < 0.67)
      throw new IOException("IO broke!")
    else throw new RuntimeException
      ("That shouldn't happen!")
  }

try {
  mightWork
} catch {
  case e:SQLException =>
```

```
    println(s"Yikes ${e.getMessage}")
  case e:IOException =>
    println(s"Uh oh ${e.getMessage}")
// this would catch unhandled exceptions
//  case e:Throwable =>
//    println(s"Groan ${e.getMessage}")
} finally {
  println("In finally block")
}
println("Finishing up.")
```

- Recall that a match block finds the first applicable match, so the more specific exceptions should be listed first
- Notice that the type Throwable applied to the last option is not required, but provides for a stylistic symmetry.
    - A warning is typically issued if this is omitted
    - If the exception isn't needed, the underscore may be used instead of an actual identifier
- If the `match` construction does not catch an exception that's thrown, then that exception propagates out of the current method (unless there's an outer `try`/`catch` construction wrapped around this one)

Iterating with `for`

- The keyword for exists for specifying iteration certain types of iteration.
- Two modes exist, the idiomatically preferred form generates a value and is called for expression.
- An alternate form provides an imperative looping construct
- Made up from up to four types of component:
    - Generators--any for construct must start with a generator. Generators create sequences of data items

from a source, such as a collection or range. A generator is of the form ***<pattern> <- <source>***
    - Filters--which remove some items from the stream. A filter is of the form **`if`** ***<boolean-expression>***
    - Variable bindings--which define/assign to an intermediate variable (of `val` type) in the body of the iteration. A binding declares and initializes a new identifier, but does not use the val prefix. It has the form ***<identifier> = <expression>***
    - A yield expression--which creates a final result value for the entire iteration. This makes a for expression, and omitting it creates an imperative for-loop. This has the form **`yield`** ***<expression>***
- A for expression must start with a generator, but may have any number of filters, variable bindings, and additional generators subsequently.
- The generators, filters, and variable binding elements follow the for keyword and are grouped using either parentheses or curly braces.
    - If parentheses are used, the elements must be separated with semicolons (because parentheses disable semicolon inferencing)
    - If curly braces are used, the elements may be placed on separate lines and semicolons are generally required.

Examples

- The following expression creates an `IndexedSequence` containing the numbers 2, 4, 6, 8, 10, 12, 14, 16, 18, 20. The boldface element is the single generator, notice that the yield keyword is followed by an e:

```
for ( x <- 1 to 10 ) yield x * 2
```

- If multiple generators are used, the later ones iterate faster than the earlier ones. This example produces a sequence of the tuples (1,1) (2,1) (2,2) (3,1) (3,2) (3,3) :

```
for (x <- 1 to 3; y <- 1 to x) yield (x, y)
```

- Filters may be added which discard elements, so the following example produce the sequence of tuples (1,2) (1,3) (2,1) (2,3) (3,1) (3,2). The boldface code is the filter (note that no parentheses are necessary on the filter construction

```
for (x <- 1 to 3;  y <- 1 to 3; if x != y) yield
(x, y)
```

- Variable bindings capture intermediate expression values which can then be used again later in the overall for expression. In this example, the variable binding is in boldface, and the code produces the output:

```
Hello starts with H
Bonjour starts with B
```

```
(for {
  w <-List("Hello", "Bonjour", "欢迎", "שָׁלוֹם")
  c1 = w.charAt(0)
  if c1 <= 'Z'
} yield s"$w starts with $c1") foreach println
```

Imperative `for` constructions

- If no yield expression is added to a for construction, an imperative structure is created.
- In this case, the for construction should be followed by a block of code, which will be executed for each data item that gets to the “end” of the sequence of elements in the for construct.
- Bindings in the for construct will be available in the block

- In this example, the boldface element is the block that's executed repeated. This prints "**`x is ?`**" with ***?*** set to each of the values 1 through 10:

```
for (x <- 1 to 10) { println(s"x is $x") }
```

- Stylistically the for expression is preferred over the imperative loop form

Expansions

- For constructions are entirely translated to an alternate representation that uses `flatmap`, `map`, and `withFilter`
- This translation only requires that the underlying types used in the generators and conditions provide those methods, without regard to implementing any particular trait or having any particular class parentage.

Function literals (lambda expressions)

Scala provides function literals (sometimes called either anonymous functions, or lambda expressions). The basic form is likely to be broadly familiar for anyone who has used this kind of feature in C++, Java, or others.

- Function literals simply define an argument list, followed by an arrow symbol, followed in turn by the expression that is the method body
    - Note that the arrow is formed with an equals sign, like JavaScript, not a minus sign like C++/Java

```
(x:Int) => x * 2
```

Several variations are possible depending on context

- If the context defines the type, the argument type might be omitted
    - Note that in this example that the form "`Int => Int`" is a type specification indicating "function taking single `Int` argument and returning `Int`

```
def doubleIt : Int => Int = (x) => x * 2
```

- If a function literal takes a single argument, without type notation, the parentheses on the argument list may be omitted

```
def doubleIt : Int => Int = x => x * 2
```

- If the types of all arguments can be reliably inferred, a shorter form is supported
- In this form, provide only the function body
    - Do not specify any argument list
    - Omit the arrow symbol
    - Arguments may be picked up using underscores
    - The underscores pick up the arguments "in order" (left to right)
    - Note that if the order of use of the arguments doesn't conveniently conform to the order in the invocation (which is determined by context) then this form is not likely to beneficial

```
def sumThree:(Int, Int, Int) => Int = _ + _ + _

def showThree:(Int, Int, Int) => String =
  "values " + _.toString + ", " +
  _.toString + ", " + _.toString
```

- A block expression may be used for the function body

```
def addAndShowThree:(Int, Int, Int) => Int =
```

```
(a, b, c) => {
  println(s"values are $a, $b, $c")
  a + b + c
}
```

- Tuple and function arguments can get confusing, it's tempting to try to use a destructuring assignment but this won't work, because the target is mistaken for a multiple argument list.
- For example, in this example, the `map` operation takes an argument that is a function that operates on the `List` elements (which are tuples of type `(Int, Int)`) and returns the sum of the two elements in each tuple. The `map` operation uses this to build a new `List` containing single Int values

```
val lt : List[(Int, Int)] = List((1, 2), (2,1))
lt map (t => t._1 + t._2) foreach println // OK

lt map((a,b) => a + b) foreach println // FAILS
```

- In the failure case, the expression `(a,b) => a + b` in the argument to `map` defines a function that takes two arguments, not a function that destructures a single tuple into two simple variables
- A destructuring can be performed using a match operation

```
lt map (_ match {case (a,b) => a + b}) foreach
println
```

- This can be shortened to

```
lt map {case (a,b) => a + b} foreach println
```

Implicit conversions, and implicit arguments

- Widening conversions (e.g. `Int` to `Long`) are automatic and as expected.
- Scala allows programmer- and library-supplied conversions too. These conversions can be made automatic in the source using implicit conversions. These conversions can be controlled/overruled by the user of the conversion.
- In addition to default arguments to functions, a function can offer a caller-configurable defaulting mechanism called implicit arguments.

Conversion of an argument

- If a function takes type Y, it may be called with a type X provided a function taking an argument of type X and returning Y is:
    - Declared in some namespace,
    - Declared as an `implicit`, and
    - In scope as a single name. E.g.

```
object DefConv {
 implicit def iToS (i: Int): String =
   "Implicitly converted " + i
 def doStuff(s: String): Unit = {
   println("doStuff with " + s)
 }
 def main(args: Array[String]): Unit = {
   val v: Int = 99;
   // Needs String, has Int, iToS converts
   doStuff(v)
```

```
  }
}
```

Conversion of the receiver

- If a method is invoked on an object (the “receiver”), but that object does not provide the method, it is possible the compiler might create a new object that does provide the desired method.
    - Receiver conversion requires an implicit function that converts the original receiver type to the type that contains the method to be invoked
    - This example illustrates a powerful use of this feature which is getting better integration of new classes with existing API
    - Note that in the example, the `+` (plus) operation invoked in the last line is the one defined in the `Count` class, not the normal “addition of `Int`” version

```
class Count(val c: Int) {
  def + (other: Count): Count =
    new Count(c + other.c)
}

implicit def intToCount(x:Int): Count =
  new Count(x)

val c1 = new Count(3)
// make Count(7), then call Count.+ method
val sum = 7 + c1
```

Implicit Arguments

- If a function argument is marked as implicit, it takes on a behavior similar to a default value, but more flexible. If, at the

point of invocation, that argument is omitted, and a value of the same type is available as a single identifier, and that identifier is labeled implicit, then that implicit value will be used as the argument to the invocation.

- Notice that this creates a situation in which the caller can decide whether to apply a default-like value to the argument, and to control what value is applied.
- It is common for library functions to also offer an implicit value for ease of use.

```
package MyLibrary {
  object ImpArg {
    object Implicits {
      implicit val defaultValue = "Hello World"
    }
    def myFunction(implicit s: String): Unit =
      println(s)
  }
}

import MyLibrary.ImpArg.myFunction
myFunction("Bonjour le monde")
import MyLibrary.ImpArg.Implicits.defaultValue
myFunction
```

- When using implicit arguments it is good practice to create a special type for the argument. In this example, having an implicit value of type String would make it very easy for the default that is implied to be applied to the wrong function.

Implicit classes

- In some situations it's desirable to create a wrapper class (that holds "utility" methods) and simultaneously define an implicit conversion to that wrapper.
    - This happens frequently when providing implicit arguments, to avoid arguments of common types like String being defined as implicit
    - It can also happen when arranging better interactions with new extension code and existing libraries.
- This situation is supported by a specific syntax called an implicit class
    - It is more restricted in application allowing conversion from only a *single* type
    - In this case the intended receiver *class* is declared as implicit
    - The implicit class must define a single-argument constructor.
    - The compiler performs the conversion using the class' constructor

```
implicit class Count(val c: Int) {
  def combine(other: Count): Count =
    new Count(c + other.c)
}

val c1 = new Count(3)
val sum = 7 combine c1
```

Implicit conversion Rules

Implicit conversions are subject to rules to avoid accidental conversions:

- Implicits cannot be top-level

- Functions must be *marked* `implicit` to be used
- A conversion must be completed with a *single* implicit
- An implicit conversion will never be used if an *explicit* conversion is coded
- The conversion must be in scope as a single identifier (this is usually achieved using imports).
    - In support of this, it's usual to place implicits in a subpackage of the package containing the affected features. This subpackage is sometimes named `implicits` or `Preamble`

Generics

- Declarations (classes, methods, variables, etc.) may all make use of Scala's generic type-consistency checking mechanism
- Declare generic type variables in square brackets.
- Place the declaration immediately after the identifier and prior to any argument list:

```
class Bucket[T](items: T*) …
def fetchAt[T](idx: Int, List[T]): T ...
```

- After declaration, generic type variables may be used in place of specific type names, and the compiler will verify consistency of usage
- Declarations may be constrained. To require that the type represented by variable T must be assignment compatible with another type X, use the syntax:

```
[T <: X]
```

- To declare that T must be assignment compatible from X:

```
[T >: X]
```

    - In either case, X may be a concrete type, or another properly declared generic type variable

- If a class declaration is intended to be covariant with respect to a generic type, the generic type declaration should be prefixed with a plus sign. If contravariance is required, prefix with a minus sign.

```
Class BucketProcessor[+T, -U]
```

- The Scala compiler will check that uses made of types are compatible with their declared variances

Functional Patterns

- Adapter/decorator: A function that takes a list of parameters of certain types may be wrapped using another function-factory that pre-applies certain arguments, or converts arguments to its own call in type, value, or both prior to passing them to the original method. Such wrapping might (depending somewhat on the particular form of declaration and use) be referred to as:
    - Pre-application
    - A curried function
    - An adapter
    - A decorator
- Monads
    - Scala collections implement the core monad behavior `flatmap`, they also provide a broad range of related and derived behaviors including: `filter`, `fold` (left and right), `map`, and `reduce` (left and right)
- `Option` is a possibly-absent value monad is provided - it has two case-class implementations `Some`, and `None`
- A monad-like class called Try allows execution of code that might throw an exception.

    - The factory for this takes a by-name argument, so that it is executed inside the Try behavior, which allows the exception to be caught:

```
Try[T](expr: => T)
```

    - The constructed `Try` object will be one of two case classes, `Success`, or `Failure`. Each of these contains the relevant data, either the result of the successful execution, or the exception that was caught.
    - Try objects provide a number of ways to get at the data they contain including: `get`, `getOrElse`, `orElse`, `recover`, `recoverWith`,
    - A Try can be converted to an `Either` (which consists of two values called left and right, left is considered unsuccessful) or to an Option object
- Streams
    - Scala Stream class provides a monadic structure that supports lazily computed infinite sequences.
    - A Stream can also be extracted from a regular collection, though in this case, the data already exist, so only the computation of intermediate values along the monad “pipeline” will be lazy.
    - Scala Stream objects are not parallel, so computations in them are handled by a single CPU
    - Apache Spark is built on top of a very similar stream concept. Spark allows execution of pipeline steps to be distributed across multiple threads on a multi-core architecture, and across distributed machines if appropriately configured.
- I/O with Scala
    - Scala has very limited IO libraries, instead it relies heavily on the core Java APIs in a JVM implementation.

- The Scala library provides a convenient means of accessing files for reading, which is the `scala.io.Source` object

Using Scala In a Maven Project

Scala can be used alongside Java in a maven project. To do this, the scala plugin can be used to compile the scala source files, and the scala library should be included in the resulting binary. The following is an example of maven pom.xml contents to achieve this.

```
<!-- Make Scala available in project.
Pay extra attention to the version -->
<dependencies>
  <dependency>
    <groupId>org.scala-lang</groupId>
    <artifactId>scala-library</artifactId>
    <version>2.12.4</version>
  </dependency>
</dependencies>

<build>
  <plugins>
    <!-- This plugin compiles Scala files -->
    <plugin>
      <groupId>net.alchim31.maven</groupId>
      <artifactId>scala-maven-plugin</artifactId>
      <version>3.3.1</version>
      <executions>
        <execution>
          <id>scala-compile-first</id>
          <phase>process-resources</phase>
          <goals>
            <goal>add-source</goal>
            <goal>compile</goal>
          </goals>
        </execution>
        <execution>
          <id>scala-test-compile</id>
          <phase>process-test-resources</phase>
          <goals>
```

```
            <goal>testCompile</goal>
          </goals>
        </execution>
      </executions>
    </plugin>
    <!-- This plugin compiles Java files -->
    <plugin>
      <groupId>org.apache.maven.plugins</groupId>
      <artifactId>maven-compiler-plugin</artifactId>
      <version>3.7.0</version>
      <configuration>
        <source>1.8</source>
        <target>1.8</target>
      </configuration>
      <executions>
        <execution>
          <phase>compile</phase>
          <goals>
            <goal>compile</goal>
          </goals>
        </execution>
      </executions>
    </plugin>
<!-- This plugin adds all dependencies to JAR file during
'package' command. Be sure to set the 'mainClass' tag.-->
    <plugin>
      <artifactId>maven-assembly-plugin</artifactId>
      <version>3.1.0</version>
      <configuration>
        <descriptorRefs>
          <descriptorRef>jar-with-dependencies</descriptorRef>
        </descriptorRefs>
        <archive>
          <manifest>
            <mainClass>REPLACE WITH MAIN CLASS FQCN</mainClass>
          </manifest>
        </archive>
      </configuration>
      <executions>
        <execution>
          <phase>package</phase>
          <goals>
            <goal>single</goal>
          </goals>
        </execution>
```

```
      </executions>
    </plugin>
  </plugins>
</build>
```